



AI for Workplace Advanced

Two-day participant guide

Build one assistant on Day 1 and keep improving it: ground it in company policy, produce on-brand documents from templates, measure its accuracy, then rebuild it as an agent workflow with human approval. Works with Copilot, ChatGPT, Claude and Gemini.

AI for Workplace (Advanced)

Two-day participant guide · Participant edition, version 0.1

© 2026 Asraf Jaafar Sidik. All rights reserved.

Online resources: github.com/Asraf-JS/AI-for-Workplace-Advanced

Microsoft, Microsoft 365, Copilot, Word and PowerPoint are trademarks of the Microsoft group of companies. ChatGPT is a trademark of OpenAI. Claude is a trademark of Anthropic. Gemini, Gemini Notebook and Google Workspace are trademarks of Google LLC. n8n is a trademark of n8n GmbH. This guide is independent training material and is not published or endorsed by any of them. All companies, people and documents in the course scenario are fictional.

Screenshots show the product at the time of writing. Microsoft updates its apps often, so labels and layouts in your tenant may differ slightly.

About the author

Asraf is a Microsoft Certified Trainer and HRD Corp-certified IT trainer based in Malaysia. He has been teaching working professionals since 2008, helping teams at banks, pharmaceutical companies, construction firms, and retailers get real work done with Microsoft 365, Power Platform, Excel, Google Workspace, and, more recently, AI tools like Copilot.

He holds a Bachelor of Information Technology (Multimedia) from the University of Malaya, along with Microsoft certifications in Power BI (PL-300) and AI for business (AB-730). Outside the classroom he writes his own automation tools in Python and Node.js, mostly to save himself from doing the same task twice.

He runs his training practice solo, which means he has personally felt every inefficiency this book tries to fix.

Contents

	Before you begin	10
1	How Large Language Models Really Behave	13
	What You Will Learn	13
	Before You Begin	13
	Topics	14
	1.1 Ask the Same Question Twice	14
	1.2 Check the Quotation with a Fast Model	14
	1.3 Check It Again with a Reasoning Model	15
	1.4 Discuss What You Saw	15
	Product Notes	16
	Independent Practice	16
	Troubleshooting	16
	Lesson Summary	16
	Hands-on exercises	17
2	Context Engineering and Structured Output	18
	What You Will Learn	18
	Before You Begin	18
	Topics	19
	2.1 Step 1: Extract the Requirements	19
	2.2 Step 2: Turn It into a Comparison Table	19
	2.3 Step 3: Turn It into JSON	20
	2.4 Compare with One Big Prompt	20
	Product Notes	20
	Independent Practice	20
	Troubleshooting	21
	Lesson Summary	21
	Hands-on exercises	22
3	Building Custom AI Assistants	24
	What You Will Learn	24
	Before You Begin	24
	Topics	25
	3.1 Create the Assistant	25
	3.2 Give It the Policy and the Vendor List	26
	3.3 Check One Quotation at a Time	26
	3.4 Compare the Three	27
	3.5 Recommend a Supplier	27
	Product Notes	28
	Independent Practice	29

	Troubleshooting	29
	Lesson Summary	29
	Hands-on exercises	30
4	Grounded Research and Source-Based Answers	32
	What You Will Learn	32
	Before You Begin	32
	Topics	33
	4.1 Set Up the Policy Notebook	33
	4.2 Ask Five Policy Questions	33
	4.3 Check Every Citation	34
	4.4 Ask Something the Policy Doesn't Cover	34
	Product Notes	35
	Independent Practice	35
	Troubleshooting	35
	Lesson Summary	35
	Hands-on exercises	37
5	Advanced Data Analysis with AI	39
	What You Will Learn	39
	Before You Begin	39
	Topics	40
	5.1 First Look at the Data	40
	5.2 Find the Data Problems	40
	5.3 Summarise the Spend	41
	5.4 Look for Patterns	42
	5.5 Make One Chart	42
	Product Notes	42
	Independent Practice	43
	Troubleshooting	43
	Lesson Summary	43
	Hands-on exercises	44
6	Controlling Output with Word and PowerPoint Templates	46
	What You Will Learn	46
	Before You Begin	46
	Topics	47
	6.1 Look Under the Hood of the Two Decks	47
	6.2 Look at the Memo's Styles	47
	6.3 Fill the Memo	47
	6.4 Make the Deck from the Good Template	48
	6.5 Do It Again with the Bad Template	49
	Product Notes	49
	Independent Practice	50

	Troubleshooting	50
	Lesson Summary	50
	Hands-on exercises	51
7	Evaluating AI Output	53
	What You Will Learn	53
	Before You Begin	53
	Topics	54
	7.1 Set Up Your Scoresheet	54
	7.2 Run the Ten Cases	54
	7.3 Score the Checker	55
	7.4 Improve One Thing and Test Again	56
	Product Notes	56
	Independent Practice	56
	Troubleshooting	56
	Lesson Summary	57
	Hands-on exercises	58
8	AI Agents, Connectors and MCP	59
	What You Will Learn	59
	Before You Begin	59
	Topics	59
	8.1 Lab: Map the systems an assistant needs to act	59
	Product Notes	59
	Independent Practice	60
	Troubleshooting	60
	Lesson Summary	60
	Hands-on exercises	61
9	Retrieval-Augmented Generation (RAG) Fundamentals	62
	What You Will Learn	62
	Before You Begin	62
	Topics	62
	9.1 Lab: Diagnose a wrong citation	62
	Product Notes	63
	Independent Practice	63
	Troubleshooting	63
	Lesson Summary	63
	Hands-on exercises	64
10	Building an AI Agent Workflow in n8n	65
	What You Will Learn	65
	Before You Begin	65
	Topics	65
	10.1 Lab: Quotation approval workflow with human approval	65

	Product Notes	66
	Independent Practice	66
	Troubleshooting	66
	Lesson Summary	66
	Hands-on exercises	67
11	AI Security for Practitioners	68
	What You Will Learn	68
	Before You Begin	68
	Topics	68
	11.1 Lab: Break-it round: hidden instructions	68
	Product Notes	68
	Independent Practice	69
	Troubleshooting	69
	Lesson Summary	69
	Hands-on exercises	70
12	AI Governance in the Malaysian Context	71
	What You Will Learn	71
	Before You Begin	71
	Topics	71
	12.1 Lab: Department AI rules and use-case register	71
	Product Notes	71
	Independent Practice	72
	Troubleshooting	72
	Lesson Summary	72
	Hands-on exercises	73
13	Capstone Project	74
	What You Will Learn	74
	Before You Begin	74
	Topics	74
	13.1 Lab: Design and pitch an AI workflow	74
	Product Notes	76
	Independent Practice	76
	Troubleshooting	76
	Lesson Summary	76
	Hands-on exercises	77
14	Product Cards	78
	The Cards	78
15	Product Card: Microsoft 365 Copilot	79
	Which Copilot Do You Have?	79
	Before Class: Check the Shield	79

Module by Module	79
Pick a Model (Module 01)	80
Upload Files	80
Build a Custom Assistant (Modules 03, 07, 10)	80
Ground Answers in Your Files (Modules 04 and 09)	81
Analyse a Spreadsheet (Module 05)	81
Word and PowerPoint (Module 06)	81
Connectors and MCP (Module 08)	81
16 Product Card: ChatGPT	82
Which Plan Do You Have?	82
Before Class: Check Your Data Setting	82
Module by Module	83
Pick a Model (Module 01)	83
Upload Files	83
Build a Custom Assistant (Modules 03, 07, 10)	84
Ground Answers in Your Files (Modules 04 and 09)	84
Analyse a Spreadsheet (Module 05)	84
Word and PowerPoint (Module 06)	84
Connectors and MCP (Module 08)	85
17 Product Card: Claude	86
Which Plan Do You Have?	86
Before Class: Check Your Data Setting	86
Module by Module	87
Turn On File Creation (Do This First)	87
Pick a Model (Module 01)	87
Upload Files	87
Build a Custom Assistant (Modules 03, 07, 10)	88
Ground Answers in Your Files (Modules 04 and 09)	88
Analyse a Spreadsheet (Module 05)	88
Word and PowerPoint (Module 06)	88
Connectors and MCP (Module 08)	88
18 Product Card: Gemini	90
Which Account Do You Have?	90
Before Class: Check Your Data Setting	90
Module by Module	91
Pick a Model (Module 01)	91
Upload Files	91
Build a Custom Assistant (Modules 03, 07, 10)	91
Ground Answers in Your Files (Modules 04 and 09)	92
Analyse a Spreadsheet (Module 05)	92
Word and PowerPoint (Module 06)	92

Connectors and MCP (Module 08)	93
--------------------------------	----

Before you begin

How to use this book

This book collects the notes and prompts from the two-day AI for Workplace (Advanced) course in one place, so you can follow along offline and keep it at your desk afterwards.

Each chapter matches one module of the course and has two parts:

- **Notes** explain each topic and walk you through the lab, with screenshots, checkpoints and troubleshooting.
- **Hands-on exercises** give you every prompt you type, ready to copy.

Every lab is written as a tool-neutral task. Each module's **Product Notes** table says what's different in Copilot Chat (Basic), M365 Copilot (Premium), ChatGPT, Claude and Gemini, and the Product Cards at the end cover the common steps in each tool.

Look out for the coloured boxes. Blue boxes are notes, teal boxes are tips, purple boxes are things to try, orange boxes are key points, and red boxes are warnings.

Prefer to copy prompts rather than retype them? Everything in this book is also online at github.com/Asraf-JS/AI-for-Workplace-Advanced, where each grey box has a copy button. Copying from there is safer than copying from a PDF, which can add line breaks in the middle of a prompt.

Your Advanced AI training journey

One assistant, built on Day 1 and improved on Day 2: grounded, tested, automated and governed.

FOUNDATIONS

01

How LLMs Behave


Why the same question gets different answers

02

Context Engineering

Prompt chains and structured output

THE PROJECT, APP BY APP





13 PITCH

Capstone

Design and pitch an AI workflow

OUTPUT

Workflow pitch

Build → Ground → Analyse → Format → Measure → Connect → Diagnose → Automate → Attack → Govern → Pitch

Course scenario

This is the continuation of **AI for Workplace (Beginner)**. You build one assistant on Day 1 and keep improving it: grounding it in company policy, producing on-brand documents from templates, measuring its accuracy, and on Day 2 rebuilding it as an agent workflow with human approval.

All lab data is for the fictional company Sinar Maju Sdn Bhd.

Ground rules

Important: use only the fictional lab data during class. Please do not upload real company, customer or personal data into any AI tool unless your organisation has approved that tool for it. AI can be wrong: check its numbers and claims before you use them.

01

CHAPTER 1

How Large Language Models Really Behave

Before you build anything, it helps to know what's going on inside the tool. An AI assistant doesn't look answers up or calculate them the way a spreadsheet does. It predicts text. That explains most of what it does well, and most of what it gets wrong. In this module you give the same supplier quotation to two kinds of model and compare what they notice.

Prompts: every prompt you need is in the steps, with a Copy button. The prompts page has them all on one page too.

Day: 1

Estimated time: 40 minutes

Your result: Two analyses of the same quotation, one from a fast model and one from a reasoning model, and a short note on how they differ.

What You Will Learn

- Explain in plain words how a large language model produces an answer
- Say why the same prompt can give a different answer each time
- Tell a fast model from a reasoning model, and when to use each
- Name the kinds of mistake to check for: arithmetic, dates and made-up details

Before You Begin

You need one AI tool open and signed in (see the product cards), and the sample files:

File	What it is
rfq-2026-118.pdf	Sinar Maju's request for quotation for 120 ergonomic mesh chairs
quotation-duduk-selesa.pdf	One supplier's reply to that request

Download both at once with [sample-files.zip](#), then right-click the ZIP and choose **Extract All**.

Note: every company, person and figure in the course files is fictional. Sinar Maju Sdn Bhd is an office supplies distributor in Petaling Jaya, and in this course you work in its Procurement Department.

Topics

Prediction, not lookup. A large language model (LLM) writes one small piece of text at a time, each time picking a likely next piece based on everything before it. It has no calculator and no database unless the tool gives it one. A total printed in a quotation is just more text to it.

Why answers vary. The model chooses among likely next words with a little randomness, so the same prompt can give a different answer each time. That's useful for drafting, and a reason to check anything that must be exact.

Fast and reasoning models. A fast model answers straight away. A reasoning (or "thinking") model works through the problem in steps before it answers, which takes longer but catches more in tasks like checking figures. Most tools let you choose.

What to check. LLMs are strong at reading, summarising and rewriting. They're weaker at arithmetic, dates, and saying "I don't know". They can also state a detail confidently that isn't in the source. Every module in this course builds a habit of checking for these.

1.1 Ask the Same Question Twice

1. Open your AI tool and start a new chat.
2. Send this prompt:

Write a one-line email subject asking a supplier to send a revised quotation. Give me only the subject line.

3. Start another new chat and send exactly the same prompt.
4. Compare the two subject lines.

The wording is almost certainly different. Neither is wrong: the model picked different likely words. Keep that in mind when you need an exact figure.

1.2 Check the Quotation with a Fast Model

1. Start a new chat and choose a fast model (the product cards show where).
2. Upload `rfq-2026-118.pdf` and `quotation-duduk-selesa.pdf`.
3. Send this prompt:

I work in procurement at Sinar Maju Sdn Bhd. I've uploaded our request for quotation (RFQ-2026-118) and one supplier's quotation. Summarise the quotation in five bullets: supplier, grand total, validity, warranty and payment terms. Then tell me in one sentence whether it meets the RFQ.

4. Write down the grand total the model gives you, and whether it says the quotation meets the RFQ.

1.3 Check It Again with a Reasoning Model

1. Start a new chat and switch to a reasoning (thinking) model.
2. Upload the same two files.
3. Send the same prompt as in 1.2, then send this one:

Now check the arithmetic. Recalculate every line amount as quantity x unit price, then the subtotal, the tax and the grand total. Show your working in a table and compare each figure with the one printed in the quotation.

4. Compare what the reasoning model found with your notes from 1.2.

Tip: don't take either model's word for a figure. Check one line yourself with a calculator: quantity times unit price.

What should you see?

The quotation prints a grand total of **RM46,548.00**, but it should be **RM47,628.00**.

Line 1 is 120 chairs at RM365.00, which is RM43,800.00. The quotation prints RM42,800.00, and the subtotal, tax and grand total all carry that RM1,000.00 gap. The amount in words matches the wrong figure, so the page looks consistent.

A fast model often repeats the printed total without checking it. A reasoning model is more likely to work out 120×365.00 and find the gap. Either can miss it, which is what you'll discuss.

1.4 Discuss What You Saw

With the person next to you, answer these three questions:

1. Did the fast model check the figures, or repeat the ones printed on the page?
2. What did the reasoning model do differently, and how much longer did it take?
3. For which tasks in your own work would you choose each kind of model?

Key point: a confident, well-formatted answer isn't a checked answer. You'll build the checking into an assistant of your own in Module 03.

Product Notes

Copilot Chat (Basic)	M365 Copilot (Premium)	ChatGPT	Claude	Gemini
Choose a quick or deeper reasoning option in the model menu at the top of the chat (it shows Auto by default). At most three files per message	Same as Basic	Choose a fast or thinking model in the model menu. Free may only offer one model	Choose a model, or turn extended thinking on, in the model menu under the message box	Choose a fast or thinking model in the model menu

Independent Practice

Give a fast model a question you already know the answer to from your own work, such as a calculation or a date difference. Then ask a reasoning model. Did either get it wrong? Which one told you how it got the answer?

Troubleshooting

Symptom	What to check
You can't find a model menu	Some free plans offer one model only. Pair up with someone on another tool or plan
The model says it can't read the PDF	Upload it again. If it still fails, open the PDF, copy all the text and paste it into the chat
Both models give the same answer	Good. Check one line with a calculator anyway: the point is to check, not to trust the slower model
The answer stops halfway	The tool is busy. Wait a minute and send the prompt again

Lesson Summary

An LLM predicts text, so it's very good with words and less reliable with sums, dates and exact facts. Reasoning models work through problems step by step and catch more, but they can still miss things. From here on, every lab includes a check you do yourself.

Check yourself: Why might a fast model repeat a wrong total printed in a quotation without noticing?

HANDS-ON

Hands-on exercises

Prompts for comparing two kinds of model on the same quotation. Hover over a grey box and click the copy icon in its top-right corner.

Part 1: The same question twice

Session: new chat, twice | **Grounding:** none

Write a one-line email subject asking a supplier to send a revised quotation. Give me only the subject line.

Part 2: Fast model

Session: new chat, fast model | **Grounding:** uploaded files (rfq-2026-118.pdf , quotation-duduk-selesa.pdf)

I work in procurement at Sinar Maju Sdn Bhd. I've uploaded our request for quotation (RFQ-2026-118) and one supplier's quotation. Summarise the quotation in five bullets: supplier, grand total, validity, warranty and payment terms. Then tell me in one sentence whether it meets the RFQ.

Part 3: Reasoning model

Session: new chat, reasoning model | **Grounding:** uploaded files (the same two)

Send the Part 2 prompt first, then:

Now check the arithmetic. Recalculate every line amount as quantity x unit price, then the subtotal, the tax and the grand total. Show your working in a table and compare each figure with the one printed in the quotation.

If the model says everything adds up:

Work out line 1 again, step by step: what is 120×365.00 ? Compare it with the amount printed on line 1.

02

CHAPTER 2

Context Engineering and Structured Output

In Module 01 you saw that a model only knows what's in front of it. Context engineering is the skill of putting the right things in front of it: the task, the background, the source, the format, and nothing that gets in the way. In this module you break one job into a chain of three short prompts, and get the output in formats a person and a machine can both use.

Prompts: every prompt you need is in the steps, with a Copy button. The prompts page has them all on one page too.

Day: 1

Estimated time: 50 minutes

Your result: A requirements checklist, a blank comparison table and a JSON version of the same requirements, all from Sinar Maju's chair RFQ.

What You Will Learn

- List what goes into a model's context: instructions, files, chat history and memory
- Write a prompt with GCSE (Goal, Context, Source, Expectations) and clear boundaries around the source
- Break a job into a chain of short prompts, each building on the last
- Ask for structured output: Markdown tables for people, JSON for other systems
- Know when to start a new chat

Before You Begin

You need `rfq-2026-118.pdf` from Module 01. If you don't have it, download [Module 01's sample-files.zip](#) and extract it.

Topics

Context is everything the model sees. That's your prompt, any files, the earlier messages in the chat, your custom instructions and anything the tool remembers about you. The model can't tell which parts matter most unless you say so. A long chat full of earlier tasks can confuse it, which is why a fresh chat often gives a better answer.

GCSE, one level up. From the Beginner course: **Goal** (what you want), **Context** (who it's for and why), **Source** (what to use) and **Expectations** (format, length, tone). At this level, add boundaries: tell the model to use only the source, and say what to do when the source doesn't answer the question.

Prompt chains. One huge prompt asks the model to do everything at once, and mistakes are hard to spot. A chain does one step per prompt. You check each result before the next step uses it.

Structured output. A Markdown table is easy for people to read and paste into Word or Excel. JSON is a format other software can read, field by field. On Day 2 your n8n workflow will need JSON.

2.1 Step 1: Extract the Requirements

1. Start a new chat in your AI tool.
2. Upload `rfq-2026-118.pdf`.
3. Send this prompt:

```
Goal: list every requirement in the attached request for quotation (RFQ-2026-118) so I can
check supplier quotations against it.
Context: I'm a procurement executive at Sinar Maju Sdn Bhd, an office supplies
distributor. We'll receive several quotations for this RFQ.
Source: use only the attached RFQ. If something isn't stated in it, don't add it.
Expectations: a Markdown table with three columns: Requirement, Exact wording from the
RFQ, How to check it in a quotation. Include the requirements for what a quotation must
contain, not only the product.
```

4. Check the table against the RFQ. Every row should come from the document, word for word in the second column.

Tip: if a row isn't in the RFQ, reply Remove any row you can't quote from the RFQ. That's the boundary in the Source line doing its job.

2.2 Step 2: Turn It into a Comparison Table

In the same chat, send:

```
Using only the requirements table above, make a blank comparison table in Markdown for three
supplier quotations. One row per requirement. Columns: Requirement, Supplier A, Supplier B,
Supplier C, Meets the RFQ? Leave the supplier cells empty. Add rows at the bottom for Grand
total (RM), Valid until and Payment terms.
```

Copy the table into Word or Excel to see that it pastes cleanly. You'll fill a table like this with real quotations in Module 03.

2.3 Step 3: Turn It into JSON

In the same chat, send:

Now give me the same requirements as JSON, so another system can read them. Use a list of objects, each with these fields: "id" (R1, R2 and so on), "requirement", "rfq_wording" and "how_to_check". Return only the JSON, with no explanation before or after it.

Check that each object matches a row of your Step 1 table. On Day 2, your n8n workflow will ask the AI for a JSON answer in the same way.

Note: the exact JSON will differ from one person to the next, as you saw in Module 01. What should match is the content: the same requirements, in fields with the names you asked for.

2.4 Compare with One Big Prompt

1. Start a **new chat** and upload `rfq-2026-118.pdf` again.
2. Send everything at once:

Read the attached RFQ, list all the requirements, make a blank comparison table for three suppliers, and give me the requirements as JSON.

3. Compare the result with your three-step chain. Which one missed a requirement, or mixed the formats up? Which would you trust to build on?

Product Notes

Copilot Chat (Basic)	M365 Copilot (Premium)	ChatGPT	Claude	Gemini
Upload the RFQ with +. Copy a table with the Copy button under the answer	Same as Basic	Tables and JSON show in formatted boxes with a copy button	Long JSON may open in a side panel; copy from there	Use Export to Sheets under a table to send it to Google Sheets

Independent Practice

Take a form or checklist from your own work (with no real data in it). Ask your AI tool to turn it into JSON with field names you choose. Then ask it to turn the JSON back into a Markdown table. Did anything get lost?

Troubleshooting

Symptom	What to check
The table has requirements that aren't in the RFQ	Ask it to quote the RFQ wording for each row, and remove rows it can't quote
The model forgot the earlier table	You started a new chat. Steps 1 to 3 must be in the same chat
The JSON has text before or after it	Ask again: "Return only the JSON."
The JSON field names are different from the ones you asked for	Paste the field names again and ask it to use exactly those

Lesson Summary

Good context is specific and bounded: the goal, the background, the one source to use, and the format you need. Short chained prompts let you check each step before the next one depends on it. And the format is part of the request: a table for people, JSON for software.

Check yourself: Why is it safer to run three short prompts in one chat than one long prompt, and when should you start a new chat instead?

HANDS-ON

Hands-on exercises

A three-step prompt chain on Sinar Maju's chair RFQ. Upload `rfq-2026-118.pdf` from [Module 01's sample files](#) before Part 1. Hover over a grey box and click the copy icon in its top-right corner.

Session guide: Parts 1 to 3 run in the **same chat**. Part 4 runs in a new chat.

Part 1: Extract the requirements

Session: new chat | **Grounding:** uploaded file (`rfq-2026-118.pdf`)

Goal: list every requirement in the attached request for quotation (RFQ-2026-118) so I can check supplier quotations against it.
Context: I'm a procurement executive at Sinar Maju Sdn Bhd, an office supplies distributor. We'll receive several quotations for this RFQ.
Source: use only the attached RFQ. If something isn't stated in it, don't add it.
Expectations: a Markdown table with three columns: Requirement, Exact wording from the RFQ, How to check it in a quotation. Include the requirements for what a quotation must contain, not only the product.

If a row isn't in the RFQ:

Remove any row you can't quote from the RFQ.

Part 2: The comparison table

Session: same chat | **Grounding:** uploaded file

Using only the requirements table above, make a blank comparison table in Markdown for three supplier quotations. One row per requirement. Columns: Requirement, Supplier A, Supplier B, Supplier C, Meets the RFQ? Leave the supplier cells empty. Add rows at the bottom for Grand total (RM), Valid until and Payment terms.

Part 3: The same requirements as JSON

Session: same chat | **Grounding:** uploaded file

Now give me the same requirements as JSON, so another system can read them. Use a list of objects, each with these fields: "id" (R1, R2 and so on), "requirement", "rfq_wording" and "how_to_check". Return only the JSON, with no explanation before or after it.

If there's text around the JSON:

Return only the JSON.

Part 4: One big prompt, for comparison

Session: new chat | **Grounding:** uploaded file (`rfq-2026-118.pdf`)

Read the attached RFQ, list all the requirements, make a blank comparison table for three suppliers, and give me the requirements as JSON.

03

CHAPTER 3

Building Custom AI Assistants

Sinar Maju has three quotations back for its chair RFQ, and more RFQs every week. Checking each quotation by hand against the policy takes time, and it's easy to miss something. In this module you build a **Quotation Checker**: a custom assistant that knows the policy and checks every quotation the same way. You'll keep improving it for the rest of the course.

Prompts: every prompt you need is in the steps, with a Copy button. The prompts page has them all on one page too.

Day: 1

Estimated time: 70 minutes

Your result: A working Quotation Checker in your AI tool, its report on each of the three chair quotations, and a recommendation you've checked yourself.

What You Will Learn

- Describe the parts of a custom assistant: instructions, knowledge, suggested prompts and tools
- Write assistant instructions with a role, steps, an output format and rules
- Build the assistant in your own AI tool
- Test it on real cases, and check its work instead of trusting it

Before You Begin

Download [sample-files.zip](#) and extract it. It holds:

File	What it is
rfq-2026-118.pdf	The request for quotation for 120 ergonomic mesh chairs
quotation-duduk-selesa.pdf	Duduk Selesa Furnishings Sdn Bhd
quotation-kerusi-nadira.pdf	Kerusi Nadira Trading
quotation-ergoluma.pdf	Ergoluma Seating Sdn Bhd
procurement-policy-v3.0.pdf	Sinar Maju's Procurement Policy, Version 3.0

`approved-vendor-list.pdf`

The suppliers Sinar Maju may buy from

Open the product card for your AI tool. It shows where to build a custom assistant in that tool.

Topics

What a custom assistant is. A saved set of instructions, and often files, that the AI follows every time you use it. Each tool has its own name for it:

Tool	Name	Can it keep files?
Copilot	Agent (Agent Builder)	Premium only. On Basic, upload files in each chat
ChatGPT	Project (personal plans) or GPT (Business and Enterprise)	Yes
Claude	Project	Yes
Gemini	Skill (personal accounts) or Gem (Workspace)	Gems yes. Skills hold plain text only, so upload PDFs in the chat

Good instructions have five parts. A **role** (who the assistant is and who it works for), a **task** (what it does, and what it doesn't), **steps** (the checks, in order), an **output format** (the same layout every time) and **rules** (what to do when it isn't sure). Instructions that only say "You are a helpful procurement assistant" leave the model to guess the rest.

Knowledge isn't a guarantee. Giving the assistant the policy means it can quote the policy. It doesn't mean it will apply every clause every time. You still test it, which is Module 07.

3.1 Create the Assistant

1. Open your AI tool and create a new custom assistant (Project, agent, Gem or skill). The product card for your tool has the clicks.
2. Name it `Quotation Checker`.
3. Paste these instructions into the instructions box:

You are the Quotation Checker for the Procurement Department of Sinar Maju Sdn Bhd, an office supplies distributor in Petaling Jaya.

Your job: check one supplier quotation against our Procurement Policy (Version 3.0), our Approved Vendor List (AVL) and the request for quotation (RFQ) it answers. You don't choose suppliers and you never correct a supplier's figures. You report what you find.

The evaluation date for every check is 10 November 2026.

Run these checks in order:

1. Arithmetic (Clause 4.4): recalculate every line amount as quantity x unit price, then any discount, the subtotal, the tax and the grand total. Compare each with the printed figure and show your working.
2. Validity (Clause 4.2): is the quotation valid for at least 30 days from its date, and

still valid on the evaluation date?

3. Supplier (Clause 5.1): is the supplier on the AVL?
4. SST (Clause 4.5): if it charges SST, does it show an SST registration number?
5. Like-for-like (Clause 4.3): does it meet every requirement in the RFQ, including specification, quantity, warranty and delivery?
6. Payment (Clause 6.2): is any deposit or advance more than 30% of the total?
7. Approval (Clause 3.1): based on the correct total including tax, who must approve?

Reply in this format:

- A table with one row per check: Check | Result (Pass or Fail) | What you found | Clause
- Correct grand total (RM), and the printed grand total if it's different
- Verdict: PASS (send for approval), REVISE (ask the supplier for a revised quotation), REJECT (supplier not on the AVL) or ESCALATE (needs the Finance Director's approval)
- One sentence explaining the verdict

Rules:

- Use only the policy, the AVL, the RFQ and the quotation. If something isn't in them, write "Not stated".
- If there's more than one problem, the verdict follows this order: REJECT, then REVISE, then ESCALATE.
- If you can't read part of a file, say so instead of guessing.

4. Save the assistant.

Note: Copilot agents allow a limited number of characters in the instructions. If yours is cut off, remove the "Reply in this format" examples last, not the checks.

3.2 Give It the Policy and the Vendor List

1. If your tool lets the assistant keep files, add `procurement-policy-v3.0.pdf` and `approved-vendor-list.pdf` to its knowledge (or project files).
2. If it doesn't (Copilot Basic, Gemini skills), you'll upload the two files at the start of each chat with the assistant.

3.3 Check One Quotation at a Time

1. Start a new chat with your Quotation Checker. If your assistant has no files, upload the policy and the vendor list first.
2. Upload `rfq-2026-118.pdf` and `quotation-duduk-selesa.pdf`, then send:

Check this quotation against RFQ-2026-118.

3. Read the report. For the arithmetic row, check one line yourself with a calculator.
4. Repeat steps 1 to 3 for `quotation-kerusi-nadira.pdf` and `quotation-ergoluma.pdf`, each in a **new chat** with the assistant.
5. Note each verdict and the clause behind it.

Tip: a new chat for each quotation stops one quotation's figures from leaking into the next report.

3.4 Compare the Three

1. Start a new chat with the assistant, and upload the RFQ and all three quotations (with Copilot, send the RFQ first, then the three quotations in a second message).
2. Send:

Build a side-by-side comparison of the three quotations, one column per supplier. Rows: quotation number, date, valid until, unit price, delivery charge, subtotal, tax, grand total, warranty, payment terms and delivery lead time. Show every figure exactly as printed in the quotation, even if it's wrong.

What should you see?

Figures as printed in each quotation:

	Duduk Selesa	Kerusi Nadira	Ergoluma
Quotation No.	DSF/QT/26/1043	KNT-2610-077	ELS/Q/2026/0388
Date	21 October 2026	22 October 2026	23 October 2026
Valid until	20 December 2026	6 December 2026	22 November 2026
Unit price	RM365.00	RM335.00	RM372.00
Delivery	RM300.00	Complimentary	RM250.00
Subtotal	RM43,100.00	RM40,200.00	RM44,890.00
Tax (SST) 8%	RM3,448.00	RM3,216.00	RM3,591.20
Grand total	RM46,548.00	RM43,416.00	RM48,481.20
Warranty	5 years frame, mechanism and gas lift	3 years frame, 1 year mechanism and gas lift	5 years frame, mechanism and gas lift
Payment	30 days	30 days	50% deposit, balance 30 days
Delivery lead time	14 to 18 days	Within 21 days	Within 21 days of deposit

Duduk Selesa's total must show **RM46,548.00**, the printed figure. If your Checker shows RM47,628.00 without saying why, it has quietly corrected the supplier's figure.

3.5 Recommend a Supplier

1. In the same chat, send:

Using your checks, recommend one supplier. Give the correct total including tax, the verdict and the clause for each supplier, what must happen before we can issue a purchase order, and who must approve it under Clause 3.1.

2. Before you accept the recommendation, check two things yourself: the arithmetic you found in 3.3, and the warranty and payment terms against the RFQ and the policy.

Show the answers

Supplier	Problem	Clause	What happens next
Duduk Selesa	Line 1 printed as RM42,800.00, but 120 x RM365.00 is RM43,800.00. Correct grand total RM47,628.00, not RM46,548.00	4.4	Ask for a revised quotation. Don't correct it yourself
Kerusi Nadira	Warranty is 3 years on the frame and 1 year on the mechanism and gas lift; the RFQ asks for 5 years on all three	4.3	Non-compliant. Ask for a revised quotation or exclude it
Ergoluma	50% deposit; the limit is 30%	6.2	Needs the Finance Director's written approval, or ask for standard terms

Recommended: **Duduk Selesa**, once it sends a revised quotation at RM47,628.00. That falls in the RM5,000 to RM50,000 band (3.1): three quotations and a comparison, approved by the HOD and the Head of Procurement.

Kerusi Nadira is cheapest, which is why a Checker that ranks on price alone gets this wrong.

Key point: keep your Quotation Checker. You'll test it on ten new quotations in Module 07, break it in Module 11 and rebuild it as a workflow in Module 10.

Product Notes

Copilot Chat (Basic)	M365 Copilot (Premium)	ChatGPT	Claude	Gemini
Agents > New agent > Skip. Instructions only: upload the policy and vendor list in each chat. At most three files per message	Same, and you can add the policy and vendor list as agent knowledge	Personal plans: a Project with instructions and files (up to 5 on Free). Business and Enterprise may build a GPT	A Project with instructions and knowledge files (up to 5 Projects on Free)	Personal: a skill (needs Keep Activity on; upload the PDFs in the chat). Workspace: a Gem with the PDFs as knowledge

Independent Practice

Think of a check you do often at work: approving a claim, reviewing a form, checking a report before it goes out. Write the five parts of an assistant for it (role, task, steps, output format, rules) without any real data. Which step would you least trust an AI to do?

Troubleshooting

Symptom	What to check
The report skips some checks	Ask: "Run all seven checks in your instructions, one row each." If it keeps skipping, check the instructions were saved in full
It says every quotation is valid	Check the evaluation date (10 November 2026) is in the instructions
It corrects the supplier's total without saying so	Ask it to show the printed figure and the correct one side by side, and remind it never to change a supplier's figures
It can't find the vendor list	Upload <code>approved-vendor-list.pdf</code> in the chat, or check it's in the assistant's knowledge
Two quotations get mixed up	Check each quotation in its own new chat

Lesson Summary

A custom assistant is instructions plus knowledge, saved so you can use them again. Clear instructions with steps, a fixed format and rules make its reports consistent, but not always right: your checker found some problems and you had to check others. That's why Module 07 tests it properly.

Check yourself: Which of the five parts of the instructions makes every report come back in the same layout, and why does that matter when you check ten quotations?

HANDS-ON

Hands-on exercises

The Quotation Checker's instructions, and the prompts to test it on the three chair quotations. Download [sample-files.zip](#) first. Hover over a grey box and click the copy icon in its top-right corner.

Part 1: The Quotation Checker instructions

Session: paste into your assistant's instructions box | **Grounding:** `procurement-policy-v3.0.pdf` and `approved-vendor-list.pdf` as knowledge, or uploaded in each chat

You are the Quotation Checker for the Procurement Department of Sinar Maju Sdn Bhd, an office supplies distributor in Petaling Jaya.

Your job: check one supplier quotation against our Procurement Policy (Version 3.0), our Approved Vendor List (AVL) and the request for quotation (RFQ) it answers. You don't choose suppliers and you never correct a supplier's figures. You report what you find.

The evaluation date for every check is 10 November 2026.

Run these checks in order:

1. Arithmetic (Clause 4.4): recalculate every line amount as quantity x unit price, then any discount, the subtotal, the tax and the grand total. Compare each with the printed figure and show your working.
2. Validity (Clause 4.2): is the quotation valid for at least 30 days from its date, and still valid on the evaluation date?
3. Supplier (Clause 5.1): is the supplier on the AVL?
4. SST (Clause 4.5): if it charges SST, does it show an SST registration number?
5. Like-for-like (Clause 4.3): does it meet every requirement in the RFQ, including specification, quantity, warranty and delivery?
6. Payment (Clause 6.2): is any deposit or advance more than 30% of the total?
7. Approval (Clause 3.1): based on the correct total including tax, who must approve?

Reply in this format:

- A table with one row per check: Check | Result (Pass or Fail) | What you found | Clause
- Correct grand total (RM), and the printed grand total if it's different
- Verdict: PASS (send for approval), REVISE (ask the supplier for a revised quotation), REJECT (supplier not on the AVL) or ESCALATE (needs the Finance Director's approval)
- One sentence explaining the verdict

Rules:

- Use only the policy, the AVL, the RFQ and the quotation. If something isn't in them, write "Not stated".
- If there's more than one problem, the verdict follows this order: REJECT, then REVISE, then ESCALATE.
- If you can't read part of a file, say so instead of guessing.

Part 2: Check one quotation

Session: new chat with the Quotation Checker, one per quotation | **Grounding:** `rfq-2026-118.pdf` and one quotation (plus the policy and vendor list if your assistant has no files)

Check this quotation against RFQ-2026-118.

If it skips a check:

Run all seven checks in your instructions, one row each.

Part 3: Compare the three

Session: new chat with the Quotation Checker | **Grounding:** the RFQ and all three quotations

Build a side-by-side comparison of the three quotations, one column per supplier. Rows: quotation number, date, valid until, unit price, delivery charge, subtotal, tax, grand total, warranty, payment terms and delivery lead time. Show every figure exactly as printed in the quotation, even if it's wrong.

Part 4: Recommend a supplier

Session: same chat | **Grounding:** the same files

Using your checks, recommend one supplier. Give the correct total including tax, the verdict and the clause for each supplier, what must happen before we can issue a purchase order, and who must approve it under Clause 3.1.

04

CHAPTER 4

Grounded Research and Source-Based Answers

Your Quotation Checker quotes the procurement policy. But when a colleague asks "how much deposit can we pay?", can you trust the answer, and can you show where it came from? In this module you set up a policy notebook that answers only from Sinar Maju's own documents, and you check every citation it gives.

Prompts: every prompt you need is in the steps, with a Copy button. The prompts page has them all on one page too.

Day: 1

Estimated time: 55 minutes

Your result: A policy notebook that answers from the Procurement Policy, five checked answers with their clause numbers, and one question it correctly says it can't answer.

What You Will Learn

- Explain the difference between an answer from the model's memory and one grounded in your sources
- Set up a source-grounded notebook in your AI tool
- Check a citation quickly, and spot one that doesn't support the answer
- Ask questions in a way that makes answers easy to check
- Recognise a good "I can't find that in the sources" answer

Before You Begin

Download [sample-files.zip](#) and extract it. It holds `procurement-policy-v3.0.pdf` and `approved-vendor-list.pdf`.

This lab works best in **Gemini Notebook** (formerly NotebookLM), which is free with a Google account and links each citation to the exact quote. Copilot Notebooks and ChatGPT or Claude Projects also work. See the Product Notes below.

Topics

Grounded and ungrounded answers. Ask a general chat "what's a typical deposit for office furniture?" and the model answers from what it learned in training: plausible, but nobody's policy. A grounded answer comes from documents you supply, and points back to them.

Source-grounded tools. A notebook (Gemini Notebook, Copilot Notebooks) or a Project keeps your sources together and answers from them. Some tools stick strictly to the sources; others mix in general knowledge unless you tell them not to.

Citations are claims too. A citation number shows where the tool says the answer came from. It can still point to the wrong clause, or to a clause that says something slightly different. Checking takes seconds if you ask for the clause number and the exact sentence.

"I don't know" is a good answer. When the sources don't cover a question, the right answer is to say so. An assistant that always has an answer is guessing some of the time.

4.1 Set Up the Policy Notebook

1. Open your notebook tool and create a new notebook (or Project) called `Sinar Maju procurement policy`.
2. Add `procurement-policy-v3.0.pdf` and `approved-vendor-list.pdf` as sources.
3. Ask:

In three bullets, what do these sources cover? Name each document and its version or date.

4. Check that both documents are named, and that the policy is shown as Version 3.0.

4.2 Ask Five Policy Questions

Send each question on its own, and keep a note of each answer and its clause number.

Answer only from the sources. For each answer, give the clause number and quote the exact sentence it comes from. What is the most we can pay in advance before the Finance Director has to approve it?

How long must a supplier's quotation be valid for?

Who has to approve a purchase of RM80,000 including tax?

Does a quotation have to show the supplier's SST registration number?

What happens if one department places several small orders with the same supplier in the same month?

Tip: the first question sets the pattern (answer only from the sources, give the clause, quote the sentence). If later answers drop the quote, say `Quote the exact sentence for that too.`

4.3 Check Every Citation

1. For each answer, open the policy PDF and find the clause it names.
2. Read the quoted sentence in the PDF. Does it say the same thing as the answer?
3. Mark each answer **Correct**, **Wrong clause** or **Doesn't support the answer**.

What should you see?

Every answer should cite the clause it comes from.

Question	Answer	Clause
What's the most you can pay in advance without extra approval?	30% of the purchase value. More needs the Finance Director's written approval first	6.2
How long must a quotation be valid?	At least 30 days from issue, and still valid on the evaluation date	4.2
Who approves a RM80,000 purchase?	The Finance Director (above RM50,000)	3.1
Must a quotation show an SST number?	Yes, if the supplier charges SST	4.5
What happens to several small orders to the same supplier?	Orders from one department to one supplier within 30 days that together cross a limit are treated as one purchase and reported to the Finance Director	7.2

If an answer has no clause number, or the clause says something different when you open it, that's the citation check failing.

4.4 Ask Something the Policy Doesn't Cover

Send:

Answer only from the sources. What is our policy on buying from suppliers outside Malaysia?

What should you see?

A good answer says the sources don't cover buying from overseas suppliers, and may point you to the Procurement Department (the policy lists procurement@sinarmaju.example). Neither the policy nor the vendor list mentions it.

If the tool gives you rules about overseas suppliers, it's answering from general knowledge, not your sources. Ask: Quote the clause that says that. It won't be able to.

Product Notes

Copilot Chat (Basic)	M365 Copilot (Premium)	ChatGPT	Claude	Gemini
Notebooks > new notebook > Add references > Upload files . Citations show as small labels with a Sources button	Same as Basic. Researcher can also search your organisation's files, with sources	A Project with the two PDFs. Ask for the clause and quote every time, because Projects also use general knowledge	A Project with the two PDFs in its knowledge. Ask for the clause and quote	Gemini Notebook (formerly NotebookLM): add both PDFs as sources. Click a citation number to see the quote

Independent Practice

Find a public document you use at work, such as a published policy or a product manual (nothing confidential). Add it to a notebook and ask three questions you already know the answers to. Check each citation. How many were exactly right?

Troubleshooting

Symptom	What to check
Answers have no clause numbers	Start each question with "Answer only from the sources" and ask for the clause and quote
The quote isn't in the PDF	The tool made it up or paraphrased. Mark it wrong, and ask it to find the exact sentence
The notebook says it can't read the PDF	Remove the source and add it again. If that fails, copy the PDF's text into a text source
Answers mention rules you can't find in the policy	The tool is mixing in general knowledge. Use a source-only tool such as Gemini Notebook, or ask it to say "not in the sources" when it can't quote one

Lesson Summary

Grounding the AI in your own documents gives answers you can trace. But a citation is only useful if someone checks it, and checking is fast when you ask for the clause and the exact sentence every time. A tool that says "that's not in the sources" is working properly.

Check yourself: An answer cites Clause 6.2, but the quoted sentence isn't in the PDF. What do you do with that answer?

HANDS-ON

Hands-on exercises

Questions for the policy notebook. Add `procurement-policy-v3.0.pdf` and `approved-vendor-list.pdf` from `sample-files.zip` as sources first. Hover over a grey box and click the copy icon in its top-right corner.

Session guide: every part runs in the **same notebook** (or Project).

Part 1: Check the sources

Session: new notebook | **Grounding:** the two PDFs as sources

In three bullets, what do these sources cover? Name each document and its version or date.

Part 2: Five policy questions

Session: same notebook | **Grounding:** sources

Answer only from the sources. For each answer, give the clause number and quote the exact sentence it comes from. What is the most we can pay in advance before the Finance Director has to approve it?

How long must a supplier's quotation be valid for?

Who has to approve a purchase of RM80,000 including tax?

Does a quotation have to show the supplier's SST registration number?

What happens if one department places several small orders with the same supplier in the same month?

If an answer drops the quote:

Quote the exact sentence for that too.

Part 3: A question the sources don't answer

Session: same notebook | **Grounding:** sources

Answer only from the sources. What is our policy on buying from suppliers outside Malaysia?

If it gives you rules anyway:

Quote the clause that says that.

05

CHAPTER 5

Advanced Data Analysis with AI

The Finance Director wants to know where Sinar Maju's money goes, and whether anyone is getting round the procurement rules. You have 21 months of purchase orders in one spreadsheet. In this module you use AI to clean the data, summarise the spend and find the patterns a person would take a day to spot, and then you check its figures.

Prompts: every prompt you need is in the steps, with a Copy button. The prompts page has them all on one page too.

Day: 1

Estimated time: 55 minutes

Your result: A cleaned view of the purchase history, a spend summary you've checked in Excel, a list of findings for the Finance Director, and one chart.

What You Will Learn

- Explain how AI tools analyse data: by running code, or by reading the text
- Find and deal with data problems before you analyse
- Define your measures so the AI calculates what you mean
- Ask for the method as well as the answer
- Check an AI figure against Excel in under a minute

Before You Begin

Download [sample-files.zip](#) and extract it. It holds the same data twice: `purchase-history.xlsx` (two sheets: the data, and **About this file**) and `purchase-history.csv`. Use the `.xlsx` unless your tool asks for CSV.

Have Excel open too. You'll use it to check one figure.

Important: the data is fictional. At work, check your organisation allows the AI tool you're using to read financial data before you upload any.

Topics

Two ways to analyse. Some tools write and run code (usually Python) on your file, then report the result: ChatGPT, Claude, Gemini and Copilot's Analyst work this way. Others read the file as text and estimate. Code is far more reliable for sums and counts. If your tool shows its code or working, it ran code.

Clean first. Real exports have duplicate rows, blanks, spelling variations and cancelled orders. Totals calculated on dirty data are wrong, however confident they look.

Say exactly what to measure. "Total spend" could mean before or after tax, with or without cancelled orders. Tell the AI which column to use and what to leave out. The **About this file** sheet tells you what each column means.

Check one figure. Pick one number and check it yourself in Excel with a filter or a PivotTable. If it matches, trust the method more. If it doesn't, ask the AI to explain the difference.

5.1 First Look at the Data

1. Start a new chat and upload `purchase-history.xlsx`.
2. Send:

This is Sinar Maju Sdn Bhd's purchase order export, January 2025 to September 2026. Read both sheets. Tell me how many rows and columns there are, the date range, and what each column means. Don't analyse anything yet.

3. Check the row count against Excel: select the PO No. column and read **Count** at the bottom of the window. It includes the header row, so subtract one.

5.2 Find the Data Problems

In the same chat, send:

Before any analysis, check the data quality. Look for: rows that are exact duplicates, rows with a blank Department, the same supplier spelled in different ways, and cancelled orders. For each problem, tell me how many rows and list the PO numbers. Don't change anything yet.

What should you see?

Before you analyse, the file needs cleaning:

- **Duplicates:** 4 rows repeat the row above exactly: PO-2025-0040, PO-2025-0192, PO-2026-0077 and PO-2026-0138.
- **Blank department:** 6 rows: PO-2025-0035, PO-2025-0153, PO-2026-0029, PO-2026-0075, PO-2026-0089 and PO-2026-0092.
- **Supplier name variants:** 7 Kertas Lestari rows are spelled "KERTAS LESTARI SDN BHD" or "Kertas Lestari Sdn. Bhd.". The Vendor ID (V001) is the same, so group by Vendor ID.

- **Cancelled orders:** 8 rows. Leave them out of spend.

5.3 Summarise the Spend

In the same chat, send:

Now clean the data: count each duplicated row once, group suppliers by Vendor ID, and leave out cancelled orders. Using Total (RM), give me: the number of unique purchase orders, total spend, spend in 2025, spend from January to September 2026, the top five suppliers by spend, and spend by category with each category's share. Show the method you used.

Then check one figure yourself:

1. In Excel, select the data and choose **Insert > PivotTable**.
2. Put **Vendor ID** in Rows and **Total (RM)** in Values.
3. Filter **Status** to leave out Cancelled.
4. Compare the total for **V003** with the AI's figure. They should match exactly.
5. Now compare **V005**. Your PivotTable is higher, because one V005 order appears twice in the file and Excel counts both. That's why you clean before you sum.

What should you see?

Using **Total (RM)**, without cancelled orders, counting each duplicate once:

Measure	Value
Purchase orders (unique)	350
Total spend	RM3,281,457.86
Spend in 2025	RM1,930,291.38
Spend January to September 2026	RM1,351,166.48

Top suppliers	Spend
Tintaria Imaging Supplies (V003)	RM966,183.56
Kertas Lestari (V001)	RM459,896.83
Pualam Paper Merchants (V002)	RM404,697.46
Duduk Selesa Furnishings (V005)	RM335,418.16
Alat Tulis Cendana (V004)	RM330,183.52

The biggest category is Stock: Toner and Ink (RM966,183.56, 29.4%), then Stock: Paper (RM857,246.33, 26.1%).

5.4 Look for Patterns

In the same chat, send:

Look for anything the Finance Director should know about. In particular:

1. Any supplier whose monthly spend changed sharply. When did it change, and by how much?
2. Orders from the same department to the same supplier within 30 days that are each below RM5,000 but add up to RM5,000 or more.
3. Orders of RM5,000 or more (Total) with fewer than 3 quotations and no Contract Ref.
4. Orders of RM5,000 or more approved only by a department head, not the Head of Procurement.
5. Any seasonal pattern in paper purchases.

For each finding, list the PO numbers and the figures.

Check at least one finding yourself: filter the spreadsheet to the PO numbers the AI lists.

Show the answers

- **Courier spend jumped in May 2026.** Kilat Merbok Express averaged RM2,388.10 a month until April 2026, then RM7,792.98 a month from May (about 3.3 times). It's under a contract, so no new quotations were needed (3.3), but it's worth asking why.
- **Split purchases.** Admin bought from Duduk Selesa three times in a week (PO-2026-0028, PO-2026-0032 and PO-2026-0036, RM14,547.60 together). IT bought from Kodbar Nusa on 18 and 19 August 2026 (PO-2026-0125 and PO-2026-0126, RM9,882.00 together). Each order was just under RM5,000 with one quotation (7.1 and 7.2).
- **Too few quotations.** PO-2025-0167 (RM8,640.00, 2 quotations) and PO-2026-0071 (RM12,409.20, 1 quotation). Both needed three (3.1).
- **Wrong approver.** PO-2026-0071 was approved by the Marketing HOD alone. It also needed the Head of Procurement.
- **Seasonal paper.** A4 paper orders in November, December and January average 4,720 reams a month, against 2,471 in other months.

5.5 Make One Chart

In the same chat, send:

Make a column chart of monthly spend with Kilat Merbok Express (V014), January 2025 to September 2026. Title it with what the chart shows, and label the axes.

Check that the chart's jump matches the month in your finding. Download the chart or take a screenshot for your notes.

Product Notes

Copilot Chat (Basic)

M365 Copilot (Premium)

ChatGPT

Claude

Gemini

Upload the `.xlsx` in a chat. Copilot reads the file, so ask it to show its calculation and check figures carefully

Analyst runs code on the file. Or open it in Excel and use Copilot there

Runs code on the file and can draw charts. Free has a data analysis limit

Turn on file creation (code execution) under **Settings > Capabilities** first. Runs code and draws charts

Upload the file in a chat. Ask it to show its working

Independent Practice

Ask one more question you'd want answered about this data, such as "Which department spends most on stationery?" Then check the answer with a PivotTable before you believe it.

Troubleshooting

Symptom	What to check
The totals don't match your PivotTable	Ask the AI which rows it included, and whether it removed duplicates and cancelled orders
It finds no duplicates	Ask it to compare complete rows, all 17 columns, not only the PO number
It can't open the <code>.xlsx</code>	Upload <code>purchase-history.csv</code> instead
It reports a pattern but no PO numbers	Ask: "List the PO numbers behind that finding." A finding you can't trace isn't a finding
It stops partway through	The analysis hit a limit. Ask for one part at a time

Lesson Summary

AI can clean, summarise and search a spreadsheet in minutes, especially when it runs code. It only calculates what you ask for, so define your measures and ask it to clean first. And one quick check in Excel tells you whether to trust the rest.

Check yourself: The AI says total spend is higher than your cleaned figure. Name two things in this file that could explain it.

HANDS-ON

Hands-on exercises

Prompts for analysing Sinar Maju's purchase history. Upload `purchase-history.xlsx` from [sample-files.zip](#) before Part 1. Hover over a grey box and click the copy icon in its top-right corner.

Session guide: every part runs in the **same chat**, the one you uploaded the file to.

Part 1: First look

Session: new chat | **Grounding:** uploaded file (`purchase-history.xlsx`)

This is Sinar Maju Sdn Bhd's purchase order export, January 2025 to September 2026. Read both sheets. Tell me how many rows and columns there are, the date range, and what each column means. Don't analyse anything yet.

Part 2: Data problems

Session: same chat | **Grounding:** uploaded file

Before any analysis, check the data quality. Look for: rows that are exact duplicates, rows with a blank Department, the same supplier spelled in different ways, and cancelled orders. For each problem, tell me how many rows and list the PO numbers. Don't change anything yet.

If it finds no duplicates:

Compare complete rows, all 17 columns, not only the PO number.

Part 3: Spend summary

Session: same chat | **Grounding:** uploaded file

Now clean the data: count each duplicated row once, group suppliers by Vendor ID, and leave out cancelled orders. Using Total (RM), give me: the number of unique purchase orders, total spend, spend in 2025, spend from January to September 2026, the top five suppliers by spend, and spend by category with each category's share. Show the method you used.

Part 4: Patterns

Session: same chat | **Grounding:** uploaded file

Look for anything the Finance Director should know about. In particular:

1. Any supplier whose monthly spend changed sharply. When did it change, and by how much?
2. Orders from the same department to the same supplier within 30 days that are each below RM5,000 but add up to RM5,000 or more.
3. Orders of RM5,000 or more (Total) with fewer than 3 quotations and no Contract Ref.
4. Orders of RM5,000 or more approved only by a department head, not the Head of Procurement.
5. Any seasonal pattern in paper purchases.

For each finding, list the PO numbers and the figures.

If a finding has no PO numbers:

List the PO numbers behind that finding.

Part 5: A chart

Session: same chat | **Grounding:** uploaded file

Make a column chart of monthly spend with Kilat Merbok Express (V014), January 2025 to September 2026. Title it with what the chart shows, and label the axes.

06

CHAPTER 6

Controlling Output with Word and PowerPoint Templates

The Quotation Checker has done its job: Duduk Selesa is the recommended supplier, once it corrects its total. Now the Head of Procurement wants an approval memo on the company template, and a short deck for the approval meeting. In this module you get AI to fill Sinar Maju's templates, and see why a well-built template matters as much as the prompt.

Prompts: every prompt you need is in the steps, with a Copy button. The prompts page has them all on one page too.

Day: 1

Estimated time: 60 minutes

Your result: An approval memo on the Sinar Maju template, a four-slide approval deck, and a side-by-side comparison of what happened with the good and the bad deck template.

What You Will Learn

- Explain what makes a Word or PowerPoint template work well with AI: styles, named layouts and placeholders
- Check a template's structure in Word and PowerPoint
- Get AI to fill a template without changing its look
- Spot the formatting problems AI output brings with it, and fix them

Before You Begin

Download [sample-files.zip](#) and extract it. It holds:

File	What it is
sinar-maju-deck.pptx	Sinar Maju's presentation template
sinar-maju-deck-bad.pptx	Another version of the same template
sinar-maju-memo.docx	Sinar Maju's approval memo template

You need desktop Word and PowerPoint, and the results from Module 03. If you don't have them, the prompts give you the facts.

Topics

Why templates. A template keeps every memo and deck on-brand without anyone thinking about fonts. It also tells an AI tool where things go, if it's built properly.

What a well-built template has. In PowerPoint: a slide master with **named layouts**, and **placeholders** for the title and content, so a new slide knows where its title goes. In Word: **styles** (Title, Heading 1, Normal) instead of text formatted by hand, and placeholder text that's clearly meant to be replaced.

What AI does with a badly built one. If every slide is a blank layout with loose text boxes, there's nothing for the AI to follow. It copies what it can see, positions drift, fonts change, and the slides have no real titles, which also makes them harder to use with a screen reader.

The finishing pass. Even with a good template, check the output: numbers, names, line breaks in tables, and anything that spills off a slide.

6.1 Look Under the Hood of the Two Decks

1. Open `sinar-maju-deck.pptx` in PowerPoint. Select **Home** > **Layout** and look at the list of layouts.
2. Select **View** > **Outline View**. Every slide's title shows in the outline.
3. Select **Review** > **Check Accessibility**.
4. Close it and do the same three things with `sinar-maju-deck-bad.pptx`.
5. Note what's different. On screen, the two decks look almost the same.

6.2 Look at the Memo's Styles

1. Open `sinar-maju-memo.docx` in Word.
2. Select **View** > **Navigation Pane**. The numbered headings are there because they use the **Heading 1** style.
3. Click into the grey text in square brackets. Those are placeholders: everything the AI (or you) must replace.

6.3 Fill the Memo

1. Start a new chat in your AI tool and upload `sinar-maju-memo.docx`.
2. Copy the prompt below, replace `[your name]` with your own name, and send it. It includes the facts from Module 03, so it works even if you don't have that chat any more.

Fill in the attached approval memo template for this purchase. Keep the template exactly as it is: the same headings, tables, styles and layout. Replace every placeholder in square brackets, and add nothing outside the template's sections.

Facts:

- To: Rozita binti Mohd Ariff, Head of Procurement. From: [your name], Procurement Executive. Date: 10 November 2026.
- Purchase: 120 ergonomic mesh office chairs for resale, RFQ-2026-118.
- Duduk Selesa Furnishings Sdn Bhd, DSF/QT/26/1043, 21 October 2026, valid until 20 December 2026. Printed total RM46,548.00, but line 1 should be RM43,800.00 (120 x RM365.00), so the correct total is RM47,628.00 (Clause 4.4). Meets the RFQ otherwise.
- Kerusi Nadira Trading, KNT-2610-077, 22 October 2026, valid until 6 December 2026, RM43,416.00. Warranty is 3 years on the frame and 1 year on the mechanism and gas lift; the RFQ asks for 5 years on all three (Clause 4.3).
- Ergoluma Seating Sdn Bhd, ELS/Q/2026/0388, 23 October 2026, valid until 22 November 2026, RM48,481.20. Asks for a 50% deposit; the limit is 30% (Clause 6.2).
- Recommendation: Duduk Selesa, subject to a revised quotation at RM47,628.00.
- Approval: RM5,000 to RM50,000 band under Clause 3.1, so the HOD and the Head of Procurement approve.

Give me the finished memo as a Word file I can download.

3. Download the file and open it in Word. Check the Navigation Pane still shows the six headings, and that no grey placeholder text is left.

If you don't see this: if your tool can't return a Word file, ask for the memo text section by section, and paste each section into the template yourself with **Paste > Keep Text Only**. That keeps the template's styles.

6.4 Make the Deck from the Good Template

1. Start a new chat and upload `sinar-maju-deck.pptx`.
2. Send:

Create a four-slide approval deck using the attached template and its layouts. Keep the template's fonts, colours and footer, and don't add any of your own.

Slide 1 (SM Title): "Approval: 120 Ergonomic Chairs for Resale", subtitle "RFQ-2026-118 recommendation", date 10 November 2026.

Slide 2 (SM Title Only, with a table): the three quotations compared: supplier, correct total (RM), warranty, payment terms, verdict.

Slide 3 (SM Content): the recommendation, Duduk Selesa at RM47,628.00 after a revised quotation, in up to four bullets.

Slide 4 (SM Closing): the decision needed from the Head of Procurement, and by when.

Remove the template's example slides. Give me the finished deck as a PowerPoint file.

3. Download it and open it in PowerPoint. Check each slide's layout (**Home > Layout**), the figures, and that nothing spills off a slide.

6.5 Do It Again with the Bad Template

1. Start a new chat, upload `sinar-maju-deck-bad.pptx`, and send the same prompt as in 6.4 (it still names the layouts, which this file doesn't have).
2. Open the result next to the deck from 6.4. Compare the titles, fonts, positions and the table.

What should you see?

The memo should fill the template without changing its look: the same headings, tables and fonts, with every grey placeholder replaced.

Field	Expected
Subject	Approval to purchase 120 ergonomic mesh office chairs for resale
Reference	RFQ-2026-118
Quotations received	Duduk Selesa (DSF/QT/26/1043), Kerusi Nadira (KNT-2610-077), Ergoluma (ELS/Q/2026/0388)
Recommended supplier	Duduk Selesa, subject to a revised quotation
Total including tax	RM47,628.00 (the corrected figure, not the printed RM46,548.00)
Approval required	HOD and Head of Procurement (Clause 3.1, RM5,000 to RM50,000)

The evaluation should name the three problems: Duduk Selesa's arithmetic (4.4), Kerusi Nadira's shorter warranty (4.3) and Ergoluma's 50% deposit (6.2).

For the deck: new slides made from `sinar-maju-deck.pptx` should use its layouts and keep its fonts and colours. Slides made from `sinar-maju-deck-bad.pptx` usually drift: titles move, fonts change, and the slides have no real titles.

Key point: the prompt was the same. The difference came from the template. If your team's AI output always needs fixing, look at the template before you rewrite the prompt.

Product Notes

Copilot Chat (Basic)

M365 Copilot (Premium)

ChatGPT

Claude

Gemini

No Copilot in Word or PowerPoint on Copilot Chat (Basic). Draft the memo in chat or a Page and paste it into the template. M365 Copilot (Basic) has standard access in the apps

Copilot in Word fills the memo in place. In PowerPoint, create a deck from the template file; work on a copy, as it can replace the open deck

Ask for a downloadable file, or use ChatGPT for Word and PowerPoint (limited on Free)

Turn on file creation under **Settings > Capabilities**. Claude returns a new .docx or .pptx

Workspace: Gemini in Docs and Slides. Personal: draft in chat and paste into the template

Independent Practice

Open a template you use at work (or one from PowerPoint's **File > New**). Check its layouts and its titles with Outline View and Check Accessibility. Would an AI tool find what it needs in it?

Troubleshooting

Symptom	What to check
The memo comes back in a different font	The tool rebuilt the document instead of filling it. Paste the text into the template with Keep Text Only
Grey placeholder text is left in the memo	Ask: "List every placeholder you didn't replace, then replace it."
The deck ignores the template's layouts	Name the layout for each slide in the prompt, as in 6.4
Text spills off a slide	Ask for fewer, shorter bullets, or move detail to the speaker notes
A figure in the deck is different from the memo	The AI retyped it. Check each figure against the facts in 6.3

Lesson Summary

AI fills a template well when the template is built well: styles in Word, named layouts and placeholders in PowerPoint. The same prompt on a badly built template gives drifting, untitled slides. Check every figure after the AI has filled the template, the same as you would a colleague's draft.

Check yourself: A colleague says "the AI keeps changing our fonts". What would you look at first in their template?

HANDS-ON

Hands-on exercises

Prompts for filling Sinar Maju's memo and deck templates. Download [sample-files.zip](#) first. Hover over a grey box and click the copy icon in its top-right corner.

Part 1: Fill the memo

Session: new chat | **Grounding:** uploaded file (`sinar-maju-memo.docx`)

Replace `[your name]` with your own name before you send it.

Fill in the attached approval memo template for this purchase. Keep the template exactly as it is: the same headings, tables, styles and layout. Replace every placeholder in square brackets, and add nothing outside the template's sections.

Facts:

- To: Rozita binti Mohd Ariff, Head of Procurement. From: `[your name]`, Procurement Executive. Date: 10 November 2026.
- Purchase: 120 ergonomic mesh office chairs for resale, RFQ-2026-118.
- Duduk Selesa Furnishings Sdn Bhd, DSF/QT/26/1043, 21 October 2026, valid until 20 December 2026. Printed total RM46,548.00, but line 1 should be RM43,800.00 (120 x RM365.00), so the correct total is RM47,628.00 (Clause 4.4). Meets the RFQ otherwise.
- Kerusi Nadira Trading, KNT-2610-077, 22 October 2026, valid until 6 December 2026, RM43,416.00. Warranty is 3 years on the frame and 1 year on the mechanism and gas lift; the RFQ asks for 5 years on all three (Clause 4.3).
- Ergoluma Seating Sdn Bhd, ELS/Q/2026/0388, 23 October 2026, valid until 22 November 2026, RM48,481.20. Asks for a 50% deposit; the limit is 30% (Clause 6.2).
- Recommendation: Duduk Selesa, subject to a revised quotation at RM47,628.00.
- Approval: RM5,000 to RM50,000 band under Clause 3.1, so the HOD and the Head of Procurement approve.

Give me the finished memo as a Word file I can download.

If placeholders are left:

List every placeholder you didn't replace, then replace it.

Part 2: The deck from the good template

Session: new chat | **Grounding:** uploaded file (`sinar-maju-deck.pptx`)

Create a four-slide approval deck using the attached template and its layouts. Keep the template's fonts, colours and footer, and don't add any of your own.

Slide 1 (SM Title): "Approval: 120 Ergonomic Chairs for Resale", subtitle "RFQ-2026-118 recommendation", date 10 November 2026.

Slide 2 (SM Title Only, with a table): the three quotations compared: supplier, correct total (RM), warranty, payment terms, verdict.

Slide 3 (SM Content): the recommendation, Duduk Selesa at RM47,628.00 after a revised quotation, in up to four bullets.

Slide 4 (SM Closing): the decision needed from the Head of Procurement, and by when.

Remove the template's example slides. Give me the finished deck as a PowerPoint file.

Part 3: The same deck from the bad template

Session: new chat | **Grounding:** uploaded file (sinar-maju-deck-bad.pptx)

Send the Part 2 prompt again, unchanged.

07

CHAPTER 7

Evaluating AI Output

Your Quotation Checker gave sensible reports on three chair quotations. Is that enough to trust it with the next fifty? In this module you test it properly: ten quotations with known answers, a score, and one improvement you can prove made it better.

Prompts: every prompt you need is in the steps, with a Copy button. The prompts page has them all on one page too.

Day: 1

Estimated time: 45 minutes

Your result: A completed scoresheet for your Quotation Checker on ten test cases, its score out of 20, and an improved version of its instructions.

What You Will Learn

- Build a test set with expected results before you test
- Score an assistant's verdicts and its reasons separately
- Tell a false positive from a false negative, and which costs more
- Change one thing at a time, and re-run every test after a change

Before You Begin

You need your **Quotation Checker** from Module 03. Download [sample-files.zip](#) and extract it. It holds:

File	What it is
<code>open-rfqs.pdf</code>	Five open requests for quotation (RFQ-2026-121 to RFQ-2026-125)
<code>case-01- ...pdf</code> to <code>case-10- ...pdf</code>	Ten supplier quotations, two for each RFQ
<code>procurement-policy-v3.0.pdf</code> , <code>approved-vendor-list.pdf</code>	The same policy and vendor list as Module 03

Some of the ten quotations comply with the policy and some don't. Don't open them to look for the problems first: the point is to see what the Checker finds on its own.

Topics

Test before you trust. "It looked right on three examples" isn't evidence. A test set is a fixed group of cases where you know the right answer before the AI sees them, so you can measure it.

Two kinds of mistake. A **false negative** passes a quotation that should fail: the costly mistake, because a bad quotation gets through. A **false positive** fails a good one: it wastes time and teaches people to ignore the Checker. Good test sets include both kinds of trap.

Score the reason, not only the verdict. "Fail" for the wrong reason is luck. Give one point for the right verdict and one for the right clause.

Change one thing at a time. When you improve the instructions, re-run all ten cases, not only the one that failed. A fix for one case can break another. That's called regression testing.

7.1 Set Up Your Scoresheet

Copy this table into Word, Excel or a notepad:

Case	Supplier	Checker verdict	Pass or Fail	Clause given	Verdict right? (1/0)	Clause right? (1/0)
01						
02						
03						
04						
05						
06						
07						
08						
09						
10						

The Checker's verdict **PASS** counts as Pass. **REVISE**, **REJECT** and **ESCALATE** all count as Fail.

7.2 Run the Ten Cases

For each case, 01 to 10:

1. Start a **new chat** with your Quotation Checker. If it has no files of its own, upload the policy and the vendor list first.
2. Upload `open-rfqqs.pdf` and the case's quotation, and send:

Check this quotation against the matching RFQ in open-rfqs.pdf. The RFQ number is in the quotation's "Your reference".

3. Fill in the Checker verdict, Pass or Fail, and the clause it gives.

Tip: working in pairs, split the cases: one person runs 01 to 05, the other 06 to 10. Then swap scoresheets to check.

7.3 Score the Checker

Open the box only once all ten rows are filled in. Give 1 point for each right verdict and 1 for each right clause.

Show the answers

Case	Supplier	Expected	Clause	Why
01	Kertas Lestari	Pass		Meets the RFQ, figures correct
02	Pualam Paper	Fail	4.4	Subtotal printed as RM17,820.00; the lines add up to RM17,280.00
03	Tintaria	Pass		The 5% discount is correct. Flagging it is a false positive
04	Dakwat Seroja	Fail	5.1	Not on the Approved Vendor List
05	Imbas Teraju	Fail	4.2	Expired on 5 November 2026
06	Kodbar Nusa	Pass		The cradle is a separate line, still like-for-like
07	Bersih Kemboja	Fail	4.5	Charges SST with no SST registration number
08	Kilau Embun	Pass		30% advance is exactly at the limit. Flagging it is a false positive
09	Rakmas	Fail	4.3	250 kg per level; the RFQ asks for 300 kg
10	Rangka Waja	Fail	4.2 and 6.2	Valid only 14 days (expired 30 October 2026) and asks for a 50% deposit

Four pass, six fail. Score 1 point for each right verdict and 1 for each right clause, out of 20. Case 10 needs both problems for both points.

Then answer:

1. How many false negatives (should fail, passed) did your Checker make? Which cases?
2. How many false positives (should pass, failed)? Which cases?
3. Which of those would cost Sinar Maju more?

7.4 Improve One Thing and Test Again

1. Pick the one mistake that matters most.
2. Open your Quotation Checker's instructions and change **one** thing. For example, add a rule:

A discount line is allowed. Recalculate the subtotal after the discount before you judge the arithmetic.

or:

An advance of exactly 30% is within the limit. Only more than 30% fails Clause 6.2.

3. Save the instructions.
4. Re-run **all ten** cases and score them again.

Did your score go up? Did any case that was right before go wrong? Keep the better version of the instructions: you'll use it in Modules 10 and 11.

Product Notes

Copilot Chat (Basic)	M365 Copilot (Premium)	ChatGPT	Claude	Gemini
Upload the policy and vendor list first, then <code>open-rfqs.pdf</code> and the case (three files per message at most). Edit the agent from Agents	Same, with the policy as agent knowledge	Run each case as a new chat inside your Project. Edit instructions in Project settings	Run each case as a new chat inside your Project. Edit the project instructions	Skill: edit it under Skills . Gem: edit the Gem's instructions

Independent Practice

Pick an AI task you rely on at work. Write five test cases where you know the right answer, including at least one false-positive trap. How would you score it?

Troubleshooting

Symptom	What to check
The Checker checks against the wrong RFQ	Tell it the RFQ number from the quotation's "Your reference"
A verdict depends on something the Checker can't see	Check the policy, the vendor list and <code>open-rfqs.pdf</code> were all in the chat or the assistant's files

Every case passes	Check the evaluation date (10 November 2026) is still in the instructions
The score went down after your change	Your change broke something else. Undo it, and try a narrower rule
You run out of time	Score the cases you finished. Five well-scored cases teach more than ten rushed ones

Lesson Summary

Testing an assistant means fixed cases, known answers and a score for both the verdict and the reason. Traps for false positives matter as much as traps for false negatives. Change one thing at a time, and re-run every case after each change.

Check yourself: Your change fixed case 03 but broke case 02. What do you do next?

HANDS-ON

Hands-on exercises

Prompts for testing your Quotation Checker on ten cases. Download [sample-files.zip](#) first. Hover over a grey box and click the copy icon in its top-right corner.

Part 1: Run a test case

Session: new chat with your Quotation Checker, one per case | **Grounding:** `open-rfqs.pdf` and one case file (plus the policy and vendor list if your assistant has no files)

Check this quotation against the matching RFQ in `open-rfqs.pdf`. The RFQ number is in the quotation's "Your reference".

Part 2: Rules you might add to the instructions

Warning: these rules give away some of the answers. Score all ten cases (7.3) before you read them.

Session: paste into your Quotation Checker's instructions, one at a time | **Grounding:** none

For a discount line flagged as an error:

A discount line is allowed. Recalculate the subtotal after the discount before you judge the arithmetic.

For an advance of exactly 30% flagged as too high:

An advance of exactly 30% is within the limit. Only more than 30% fails Clause 6.2.

For a missing SST number that wasn't flagged:

If a quotation has a tax or SST line, look for an SST registration number in the letterhead and the footer. If there isn't one, Clause 4.5 fails.

08

CHAPTER 8

AI Agents, Connectors and MCP

One short paragraph on where this module sits in the course, to be added.

Prompts: every prompt you need is in the steps, with a Copy button. The prompts page has them all on one page too.

Day: 2

Estimated time: TBA

Your result: To be added

What You Will Learn

To be added

Before You Begin

To be added

Topics

To be added

8.1 Lab: Map the systems an assistant needs to act

Steps to be added.

Product Notes

Copilot Chat (Basic)	M365 Copilot (Premium)	ChatGPT	Claude	Gemini
TBA	TBA	TBA	TBA	TBA

Independent Practice

To be added

Troubleshooting

Symptom	What to check
TBA	TBA

Lesson Summary

To be added

HANDS-ON

Hands-on exercises

The prompts for this module's lab. Hover over a grey box and click the copy icon in its top-right corner.

Part 1: *To be added*

Session: *TBA* | **Grounding:** *TBA*

Prompt to be added.

09

CHAPTER 9

Retrieval-Augmented Generation (RAG) Fundamentals

One short paragraph on where this module sits in the course, to be added.

Prompts: every prompt you need is in the steps, with a Copy button. The prompts page has them all on one page too.

Day: 2

Estimated time: TBA

Your result: To be added

What You Will Learn

To be added

Before You Begin

To be added

Topics

To be added

9.1 Lab: Diagnose a wrong citation

Steps to be added.

Show the answers

The wrong answers come from Version 2.1, which says SUPERSEDED only on its first page.

Question	Right answer (v3.0)	Wrong answer (from v2.1)	Clause
----------	---------------------	--------------------------	--------

Most you can pay in advance without extra approval?	30%	50%	6.2
How long must a quotation be valid?	At least 30 days	At least 14 days	4.2
Who approves a RM80,000 purchase?	Finance Director	HOD and Head of Procurement	3.1
Must a quotation show an SST number?	Yes, if SST is charged	A clause about the lowest price	4.5

The clause numbers are the same in both versions, which is why a wrong citation looks right. Open the cited page and check the version in its header.

Product Notes

Copilot Chat (Basic)	M365 Copilot (Premium)	ChatGPT	Claude	Gemini
TBA	TBA	TBA	TBA	TBA

Independent Practice

To be added

Troubleshooting

Symptom	What to check
TBA	TBA

Lesson Summary

To be added

HANDS-ON

Hands-on exercises

The prompts for this module's lab. Hover over a grey box and click the copy icon in its top-right corner.

Part 1: *To be added*

Session: *TBA* | **Grounding:** *TBA*

Prompt to be added.

10

CHAPTER 10

Building an AI Agent Workflow in n8n

One short paragraph on where this module sits in the course, to be added.

Prompts: every prompt you need is in the steps, with a Copy button. The prompts page has them all on one page too.

Day: 2

Estimated time: TBA

Your result: To be added

What You Will Learn

To be added

Before You Begin

To be added

Topics

To be added

10.1 Lab: Quotation approval workflow with human approval

Steps to be added.

Show the answers

Each quotation should take a different route through your workflow:

Quotation	Route	Clause	Why
Alat Tulis Cendana	Send for approval	3.1	Meets the RFQ and the policy. RM39,852.00, so the HOD and the Head of Procurement approve

Kodbar Nusa	Ask for a revised quotation	4.4	Subtotal printed as RM36,800.00; the lines add up to RM36,080.00. Correct total RM38,966.40, not RM39,744.00
Hancur Rapi Supplies	Reject and start supplier registration	5.1	Not on the Approved Vendor List, though it's the cheapest
Imbas Teraju	Escalate to the Finance Director	6.2	Asks for a 40% deposit; the limit is 30%

Nothing should reach a supplier, or be approved, until a person clicks Approve in the Send and Wait step.

Product Notes

Copilot Chat (Basic)	M365 Copilot (Premium)	ChatGPT	Claude	Gemini
TBA	TBA	TBA	TBA	TBA

Independent Practice

To be added

Troubleshooting

Symptom	What to check
TBA	TBA

Lesson Summary

To be added

HANDS-ON

Hands-on exercises

The prompts for this module's lab. Hover over a grey box and click the copy icon in its top-right corner.

Part 1: *To be added*

Session: *TBA* | **Grounding:** *TBA*

Prompt to be added.



CHAPTER 11

AI Security for Practitioners

One short paragraph on where this module sits in the course, to be added.

Prompts: every prompt you need is in the steps, with a Copy button. The prompts page has them all on one page too.

Day: 2

Estimated time: TBA

Your result: *To be added*

What You Will Learn

To be added

Before You Begin

To be added

Topics

To be added

11.1 Lab: Break-it round: hidden instructions

Steps to be added.

Product Notes

Copilot Chat (Basic)	M365 Copilot (Premium)	ChatGPT	Claude	Gemini
TBA	TBA	TBA	TBA	TBA

Independent Practice

To be added

Troubleshooting

Symptom	What to check
TBA	TBA

Lesson Summary

To be added

HANDS-ON

Hands-on exercises

The prompts for this module's lab. Hover over a grey box and click the copy icon in its top-right corner.

Part 1: *To be added*

Session: *TBA* | **Grounding:** *TBA*

Prompt to be added.

12

CHAPTER 12

AI Governance in the Malaysian Context

One short paragraph on where this module sits in the course, to be added.

Prompts: every prompt you need is in the steps, with a Copy button. The prompts page has them all on one page too.

Day: 2

Estimated time: TBA

Your result: To be added

What You Will Learn

To be added

Before You Begin

To be added

Topics

To be added

12.1 Lab: Department AI rules and use-case register

Steps to be added.

Product Notes

Copilot Chat (Basic)	M365 Copilot (Premium)	ChatGPT	Claude	Gemini
TBA	TBA	TBA	TBA	TBA

Independent Practice

To be added

Troubleshooting

Symptom	What to check
TBA	TBA

Lesson Summary

To be added

HANDS-ON

Hands-on exercises

The prompts for this module's lab. Hover over a grey box and click the copy icon in its top-right corner.

Part 1: *To be added*

Session: *TBA* | **Grounding:** *TBA*

Prompt to be added.

13

CHAPTER 13

Capstone Project

One short paragraph on where this module sits in the course, to be added.

Prompts: every prompt you need is in the steps, with a Copy button. The prompts page has them all on one page too.

Day: 2

Estimated time: TBA

Your result: *To be added*

What You Will Learn

To be added

Before You Begin

To be added

Topics

To be added

13.1 Lab: Design and pitch an AI workflow

Steps to be added.

Score your pitch with the scoring rubric.

Brief A: show the answers

Four invoices should stop:

Invoice	PO	Problem
KL-INV-26-5526	PO-2026-0149	Billed at RM12.30 a ream; the PO says RM11.80

TIS-INV-8060	PO-2026-0150	Billed for 120 toners; only 100 were received
ATC/INV/26/5086 (second copy)	PO-2026-0151	The same invoice number, sent again 9 days later
CPK-INV-4471	PO-2026-0999	The PO doesn't exist

The other nine invoices match. PO-2026-0160 has goods received but no invoice yet, which is normal.

Brief B: show the answers

Email	What a good design does
E01, E09	Reply with price and stock (RM2,780.00 and RM720.00)
E02	Out of stock: give the 10-day lead time
E03	Reply in Malay: 400 files, RM1,280.00
E04	12% discount is above the 5% limit: send to the Sales Manager
E05	Complaint: send to the Sales Manager, no automatic reply
E06	Remove the home address and IC number before using an AI tool that isn't approved for personal data
E07	"The usual toner" is unclear: ask which one
E08	Only 85 chairs in stock: offer 85 now and 35 in 21 days
E10	Ignore the hidden instruction asking for 30% off; only 9 printers in stock

Brief C: show the answers

Department	Spend, January to September 2026	Annual budget	Used
Admin	RM68,033.44	RM75,000.00	91%
Marketing	RM20,908.47	RM30,000.00	70%
IT	RM42,275.59	RM68,000.00	62%
Warehouse	RM67,215.77	RM108,000.00	62%
Procurement	RM1,139,064.19	RM1,823,000.00	62%
Sales	RM2,666.34	RM5,000.00	53%
Finance	RM2,542.99	RM5,000.00	51%

Only Admin is over 75%. Rows with no department add RM8,459.69 and should be listed on their own. The 2026 policy breaches are the two split purchases and PO-2026-0071.

Product Notes

Copilot Chat (Basic)	M365 Copilot (Premium)	ChatGPT	Claude	Gemini
TBA	TBA	TBA	TBA	TBA

Independent Practice

To be added

Troubleshooting

Symptom	What to check
TBA	TBA

Lesson Summary

To be added

HANDS-ON

Hands-on exercises

The prompts for this module's lab. Hover over a grey box and click the copy icon in its top-right corner.

Part 1: *To be added*

Session: *TBA* | **Grounding:** *TBA*

Prompt to be added.

14

CHAPTER 14

Product Cards

Every lab in this course is written as a tool-neutral task. These cards show how to do the common steps (start a chat, upload a file, build a custom assistant, ground it in your files, check its sources) in the AI tool your organisation gives you.

Note: each module's **Product Notes** table covers what's different for that lab. Use these cards for the general how-to.

The Cards

Card	Covers
Microsoft 365 Copilot	Copilot Chat (Basic) and M365 Copilot (Premium)
ChatGPT	Free and paid plans
Claude	Free and paid plans
Gemini	Free and Google Workspace

Note: the cards were checked in October 2026. These tools change every month, so a menu may have moved by the time you read this. Tell your trainer if something doesn't match your screen.

15

CHAPTER 15

Product Card: Microsoft 365 Copilot

How to do this course's labs in Microsoft 365 Copilot, with or without the paid license. Use it alongside each module's **Product Notes** table.

Note: checked in October 2026. Copilot changes often, so a button may have moved. If something here doesn't match your screen, tell your trainer.

Which Copilot Do You Have?

Open the **Microsoft 365 Copilot** app (m365.cloud.microsoft) with your work account. The label under your name at the bottom left tells you which Copilot you have.

Label	Who has it	What matters for this course
Copilot Chat (Basic)	Organisations with more than 2,000 users, without the paid license	Chat, file upload, Notebooks, Pages and simple agents. No Copilot inside Word, Excel or PowerPoint
M365 Copilot (Basic)	Smaller organisations, without the paid license	As above, plus standard access to Copilot in Word, Excel and PowerPoint
M365 Copilot (Premium)	People with the paid Microsoft 365 Copilot license	Everything above, plus your work data, Researcher, Analyst, and full Copilot in the Office apps

Before Class: Check the Shield

Look for the green shield at the top right of the chat. Hover over it: it should say **Enterprise data protection applies to this chat**. With the shield, your prompts and files are protected under your organisation's Microsoft agreement and aren't used to train the models.

! Stop: if you don't see the shield, you may be signed in with a personal Microsoft account. Sign out and sign in with your work account before you upload anything.

Module by Module

Module	What you do	Basic	Premium
--------	-------------	-------	---------

01	Compare two model types	Switch models at the top of the chat	Same
02	Prompt chain with Markdown output	Any chat	Same
03	Build the Quotation Checker	Agent Builder: instructions only, so upload the policy in each chat	Agent Builder, with the policy as agent knowledge
04	Grounded research with citations	A Notebook with the policy files	Same, or Researcher
05	Analyse the purchase history	Upload the spreadsheet in a chat	Analyst , or Copilot in Excel
06	Memo and deck from templates	Draft in chat or a Page, then paste into the template	Copilot in Word and PowerPoint, from the template
07	Score the Checker	Run each test quotation with your agent	Same
08	Connectors and MCP	Look at what your agents can reach	Copilot connectors and Copilot Studio (admin-managed)
09	Diagnose a wrong citation	A Notebook with both policy versions	Same

Pick a Model (Module 01)

1. Start a new chat.
2. At the top of the chat, open the model menu (it shows **Auto** by default).
3. Send the same prompt with a quick option, then with a deeper reasoning option.

Upload Files

Click **+** in the message box and choose your files.

If you don't see this: Copilot takes at most three files per message. For labs with more files, send the policy first in its own message, then the quotations.

Build a Custom Assistant (Modules 03, 07, 10)

1. In the left pane, select **Agents**, then **New agent**.
2. Select **Skip** to go straight to the form.
3. Name it `Quotation Checker` and paste the Checker instructions from the Module 03 prompts page into **Instructions**.
4. Add one or two **Suggested prompts**, such as "Check this quotation".
5. Select **Create**, then test it.

On **Basic**, agents can use instructions and public websites only, so upload the policy and vendor list in each chat with the agent. On **Premium**, you can also add files or SharePoint as the agent's knowledge.

If you don't see this: your admin may have turned off Agent Builder. Paste the instructions at the start of a normal chat instead.

Ground Answers in Your Files (Modules 04 and 09)

1. In the left pane, select **Notebooks**, then create a new notebook.
2. Select **Add references**, then **Upload files**, and add the policy and the vendor list.
3. Ask your question, and ask for the clause number with every answer.
4. Open the PDF and check the clause says what Copilot claims.

On **Premium**, **Researcher** can also search your organisation's files and the web, with sources.

Analyse a Spreadsheet (Module 05)

- **Basic:** upload `purchase-history.xlsx` in a chat and ask your question. Ask Copilot to show its calculation, and check one figure yourself in Excel.
- **Premium:** use **Analyst**, which runs code on your data, or open the file in Excel and use Copilot there.

If you don't see this: Analyst and Researcher may appear in the left pane on a Basic account, but they need the paid license.

Word and PowerPoint (Module 06)

- **Basic:** draft the memo in chat, select **Edit in Pages**, and export it with **More actions** > **Export** > **Document**. Then copy the text into `sinar-maju-memo.docx`.
- **Premium:** open `sinar-maju-memo.docx` in Word and use Copilot there. In PowerPoint, Copilot can create a deck from a file or template. It follows a template best when the layouts have clear names, which is the difference between `sinar-maju-deck.pptx` and `sinar-maju-deck-bad.pptx`.

! Caution: in PowerPoint, creating a presentation from a file can replace the open presentation. Work on a copy.

Connectors and MCP (Module 08)

Copilot reaches other systems through **Copilot connectors** and through agents built in **Copilot Studio**, which can use MCP servers. Your IT admin sets these up, not you. Module 08 is a mapping exercise, so you don't need to connect anything.

16

CHAPTER 16

Product Card: ChatGPT

How to do this course's labs in ChatGPT. Use it alongside each module's **Product Notes** table.

Note: checked in October 2026. ChatGPT changes often, so a button may have moved. If something here doesn't match your screen, tell your trainer.

Which Plan Do You Have?

Click your name or picture at the bottom left. Your plan shows in the menu.

Plan	Who has it	What matters for this course
Free	Anyone with an account	Works for most labs, with tight daily limits on uploads and data analysis
Go, Plus, Pro	Personal paid plans	Higher limits. Same features as Free for most labs
Business, Enterprise, Edu	A company or school workspace	Your data isn't used for training by default, and you may be able to build GPTs

Before Class: Check Your Data Setting

On personal plans (Free, Go, Plus, Pro), ChatGPT can use your chats to train its models unless you switch that off.

1. Click your name or picture, then **Settings**.
2. Select **Data controls**.
3. Turn off **Improve the model for everyone**.

Important: even with that setting off, use only the course's fictional files. Don't upload real company documents to a personal plan.

On Business, Enterprise and Edu, your chats aren't used for training by default. Your workspace admin controls the setting.

Module by Module

Module	What you do	In ChatGPT
01	Compare two model types	Switch models in the model menu
02	Prompt chain with Markdown output	Any chat
03	Build the Quotation Checker	A Project with instructions and files (personal plans can't build new GPTs)
04	Grounded research with citations	A Project with the policy files, or use Gemini Notebook for this lab
05	Analyse the purchase history	Upload the spreadsheet; ChatGPT runs the analysis
06	Memo and deck from templates	Upload the template and ask for a file, or use ChatGPT for Word and PowerPoint
07	Score the Checker	Run each test quotation in your Project
08	Connectors and MCP	Look at Apps in Settings
09	Diagnose a wrong citation	A Project with both policy versions

Pick a Model (Module 01)

1. Start a new chat.
2. Open the model menu at the top of the chat.
3. Choose a fast model for one answer, then a thinking (reasoning) model for the second.

If you don't see this: on Free, you may only have one model, or ChatGPT may switch you to a smaller one after a few messages. Compare with someone on another plan or tool instead.

Upload Files

1. In the message box, click +.
2. Choose your file and send your prompt.

If you don't see this: Free plans allow only a few uploads a day, and fewer at busy times. Upload only the files the lab needs, or put several into one Project (see below) so you don't upload them again.

Build a Custom Assistant (Modules 03, 07, 10)

Personal ChatGPT plans can't create new GPTs any more, so the Quotation Checker is built as a **Project**. A Project keeps your instructions, files and chats together, and every chat inside it follows the instructions.

1. In the left pane, next to **Projects**, select **New project**.
2. Name it `Quotation Checker` and create it.
3. Open the project's **more options** menu (...) and choose **Project settings**.
4. Paste the Checker instructions from the Module 03 prompts page into the instructions box, and save.
5. Add the policy (`procurement-policy-v3.0.pdf`) and the vendor list (`approved-vendor-list.pdf`) as project files.
6. Start a chat inside the project and upload a quotation to check it.

Tip: Free plans allow up to 5 files per project, which is enough for the policy, the vendor list and a quotation or two.

On a **Business, Enterprise or Edu** workspace you may be able to build a GPT instead, if your admin allows it.

Ground Answers in Your Files (Modules 04 and 09)

Use a Project with only the policy files in it, and ask for the clause number with every answer. Then open the PDF and check the clause says what ChatGPT claims.

For Module 04, many people find **Gemini Notebook** (free with a Google account) easier, because each citation links to the exact quote. See the Gemini card.

Analyse a Spreadsheet (Module 05)

1. Upload `purchase-history.xlsx`.
2. Ask your question. ChatGPT writes and runs code to work it out, and can make charts.
3. Ask it to show the calculation, so you can check it.

If you don't see this: data analysis has its own limit on Free. If you hit it, use the `.csv` file, ask fewer, bigger questions, or pair up with someone on a paid plan.

Word and PowerPoint (Module 06)

Two ways:

- **In the chat:** upload `sinar-maju-memo.docx` or `sinar-maju-deck.pptx`, describe what you need, and ask for a finished file to download.

- **Inside Office: ChatGPT for Word** and **ChatGPT for PowerPoint** add a ChatGPT pane to the app. They're on every plan, with limited use on Free.

Connectors and MCP (Module 08)

Click your name, then **Settings**, then **Apps** to see the apps ChatGPT can connect to, such as Google Drive or Outlook. Adding your own MCP server needs developer mode, which is fully available only on Business, Enterprise and Edu. Module 08 is a mapping exercise, so you don't need to connect anything.

17

CHAPTER 17

Product Card: Claude

How to do this course's labs in Claude. Use it alongside each module's **Product Notes** table.

Note: checked in October 2026. Claude changes often, so a button may have moved. If something here doesn't match your screen, tell your trainer.

Which Plan Do You Have?

Click your name at the bottom left. Your plan shows next to it.

Plan	Who has it	What matters for this course
Free	Anyone with an account	Works for most labs, with lower usage limits. Up to 5 Projects
Pro, Max	Personal paid plans	Higher limits, more models, and the Claude for Word, Excel and PowerPoint add-ins
Team, Enterprise	A company workspace	Your admin controls settings such as file creation and connectors

Before Class: Check Your Data Setting

On Free, Pro and Max, you choose whether your chats can be used to train Claude.

1. Click your name, then **Settings**.
2. Select **Privacy**.
3. Turn off the setting that lets your chats help improve Claude's models.

Important: even with that setting off, use only the course's fictional files. Don't upload real company documents to a personal plan.

On Team and Enterprise, chats aren't used for training.

Module by Module

Module	What you do	In Claude
01	Compare two model types	Switch models, or turn extended thinking on and off
02	Prompt chain with Markdown output	Any chat
03	Build the Quotation Checker	A Project with instructions and knowledge files
04	Grounded research with citations	A Project with the policy files, or use Gemini Notebook for this lab
05	Analyse the purchase history	Upload the spreadsheet; Claude runs the analysis
06	Memo and deck from templates	Upload the template and ask Claude to create the file
07	Score the Checker	Run each test quotation in your Project
08	Connectors and MCP	Customize > Connectors
09	Diagnose a wrong citation	A Project with both policy versions

Turn On File Creation (Do This First)

Modules 05 and 06 need Claude to run code and create files.

1. Click your name, then **Settings**.
2. Select **Capabilities**.
3. Turn on code execution and file creation.

On Team and Enterprise you won't see this switch. Your organisation's owner controls it.

Pick a Model (Module 01)

1. Start a new chat.
2. Open the model menu under the message box.
3. Send the same prompt once with a fast model and once with extended thinking on (or a larger model).

Upload Files

Click **+** in the message box, choose **Upload a file**, and send your prompt. You can also drag files into the chat.

Build a Custom Assistant (Modules 03, 07, 10)

A **Project** keeps your instructions and files together, and every chat inside it follows the instructions. Free accounts can have up to 5 Projects.

1. In the left pane, select **Projects**, then **New project**.
2. Name it `Quotation Checker` and create it.
3. Under the project's instructions, paste the Checker instructions from the Module 03 prompts page, and save.
4. Add `procurement-policy-v3.0.pdf` and `approved-vendor-list.pdf` to the project knowledge.
5. Start a chat inside the project and upload a quotation to check it.

Ground Answers in Your Files (Modules 04 and 09)

Use a Project with only the policy files in it, and ask for the clause number and a short quote with every answer. Then open the PDF and check the quote is really there.

For Module 04, many people find **Gemini Notebook** (free with a Google account) easier, because each citation links to the exact quote. See the Gemini card.

Analyse a Spreadsheet (Module 05)

1. Check that file creation is on (see above).
2. Upload `purchase-history.xlsx`.
3. Ask your question. Claude writes and runs code to work it out, and can make charts.
4. Ask it to show the calculation, so you can check it.

Word and PowerPoint (Module 06)

Claude can create Word (.docx), PowerPoint (.pptx), Excel (.xlsx) and PDF files on every plan, once file creation is on.

1. Upload `sinar-maju-memo.docx` or `sinar-maju-deck.pptx`.
2. Describe what you need and ask Claude to fill the template and give you the file.
3. Download the file and open it in Word or PowerPoint.

Claude makes a new file rather than changing your original. On Pro, Max, Team and Enterprise, the **Claude for Word**, **Claude for PowerPoint** and **Claude for Excel** add-ins can also edit a document directly inside the app.

Connectors and MCP (Module 08)

1. Click **Customize** in the left pane, then **Connectors**.

2. Browse the directory to see which tools Claude can connect to, and what each one can read or change.

Claude connects to MCP servers as custom connectors (**Add custom connector**). Free accounts can add one; paid plans can add more. Module 08 is a mapping exercise, so you don't need to connect anything.

 **Caution:** a custom connector isn't reviewed by Anthropic. Only add one your organisation trusts.

18

CHAPTER 18

Product Card: Gemini

How to do this course's labs in Gemini and Gemini Notebook. Use it alongside each module's **Product Notes** table.

Note: checked in October 2026. Google is changing Gemini quickly: Gems are being replaced by skills, and NotebookLM was renamed **Gemini Notebook** in July 2026. If something here doesn't match your screen, tell your trainer.

Which Account Do You Have?

This matters more in Gemini than in the other tools, because the account type decides both the features and how your data is treated.

Account	Who has it	What matters for this course
Personal Google account (free)	Anyone with Gmail	Most labs work. Deep Research is limited. Skills need you to be 18 or over
Google AI Plus, Pro or Ultra	Personal paid plans	Higher limits, including in Gemini Notebook
Google Workspace (work or school)	Your company or school	Licensed accounts: chats aren't reviewed by people or used for training. Gems stay until March 2027

To check: open gemini.google.com and look at the account picture at the top right. A work or school account shows your company's domain.

Before Class: Check Your Data Setting

On a **personal** account, Gemini keeps your chats and may use them to train its models, with some chats read by human reviewers, while **Keep Activity** is on.

1. In Gemini, open **Settings and help**, then **Activity**.
2. Turn **Keep Activity** off if you don't want your chats used for training.

Important: skills (below) only work with Keep Activity **on**. If you turn it on to use a skill, use only the course's fictional files, which you should do anyway.

On a licensed **Workspace** account, chats and uploaded files aren't reviewed by people or used to train models.

Module by Module

Module	What you do	In Gemini
01	Compare two model types	Switch models in the model menu
02	Prompt chain with Markdown output	Any chat
03	Build the Quotation Checker	A skill (personal accounts) or a Gem (Workspace, until March 2027)
04	Grounded research with citations	Gemini Notebook (formerly NotebookLM)
05	Analyse the purchase history	Upload the spreadsheet in a chat
06	Memo and deck from templates	Gemini in Docs and Slides (Workspace), or draft in a chat and paste into the template
07	Score the Checker	Run each test quotation with your skill or Gem
08	Connectors and MCP	Look at the apps Gemini can connect to
09	Diagnose a wrong citation	A Gemini Notebook with both policy versions

Pick a Model (Module 01)

1. Start a new chat.
2. Open the model menu (next to the message box or at the top).
3. Send the same prompt with a fast model, then with a thinking model.

Upload Files

Click **+** in the message box, choose **Upload files**, and send your prompt.

Build a Custom Assistant (Modules 03, 07, 10)

Personal account: a skill

Google is replacing Gems with **skills**, and Gems on personal accounts end in November 2026. A skill is a saved set of instructions Gemini uses when it fits the task.

1. Make sure **Keep Activity** is on (see above).
2. Open **Settings and help**, then **Skills**.
3. Choose **Create manually**.

4. Name it `quotation-checker`, add a one-line description starting "Use when checking a supplier quotation", and paste the Checker instructions from the Module 03 prompts page.
5. Select **Create**.
6. In a new chat, upload the policy, the vendor list and a quotation, and ask Gemini to check the quotation.

If you don't see this: skills need a personal Google account and an age of 18 or over. On a work or school account, use a Gem instead.

Skills can't hold PDF, Word or Excel files as references, only plain text. That's why you upload the policy and vendor list in the chat.

Workspace account: a Gem

1. In the left pane, select **Gems**, then **New Gem**.
2. Name it `Quotation Checker` and paste the Checker instructions.
3. Under **Knowledge**, upload `procurement-policy-v3.0.pdf` and `approved-vendor-list.pdf`.
4. Save, then start a chat with the Gem and upload a quotation.

Gems on Workspace business accounts end in March 2027 and move to skills automatically.

Ground Answers in Your Files (Modules 04 and 09)

Gemini Notebook (formerly NotebookLM) answers only from the sources you add, and each citation shows the exact quote it used.

1. Go to notebooklm.google and sign in.
2. Create a new notebook.
3. Add `procurement-policy-v3.0.pdf` and `approved-vendor-list.pdf` as sources.
4. Ask your question, then click a citation number to see the quote in the source.

Free accounts can add up to 50 sources to a notebook, and usage limits reset every five hours.

Analyse a Spreadsheet (Module 05)

1. Upload `purchase-history.xlsx` (or the `.csv`) in a chat.
2. Ask your question, and ask Gemini to show how it worked it out.
3. Check one figure yourself in Excel or Google Sheets.

Word and PowerPoint (Module 06)

- **Workspace account with Gemini:** open the memo template in Google Docs or the deck in Google Slides, and use the Gemini side panel to draft into it.
- **Personal account:** draft the memo text in a chat, then paste it into `sinar-maju-memo.docx` yourself. Gemini works best with Google's own formats, so expect to tidy the formatting.

Connectors and MCP (Module 08)

On a personal account, Gemini can connect to Google apps such as Gmail and Drive. On Workspace, your admin decides which apps it can reach. Module 08 is a mapping exercise, so you don't need to connect anything.

! **Caution:** when you connect Gmail or Drive on a personal account, parts of what Gemini reads there can be used to improve Google's models. Don't connect a work mailbox to a personal account.

AI for Workplace (Advanced)

This book accompanies the two-day AI for Workplace (Advanced) course, the continuation of AI for Workplace (Beginner). Participants build a Procurement Quotation Checker for the fictional company Sinar Maju Sdn Bhd and keep improving it: grounding it in company policy, producing on-brand documents from templates, measuring its accuracy, and rebuilding it as an agent workflow with human approval, before covering AI security and governance in the Malaysian context. Every lab works with Copilot, ChatGPT, Claude or Gemini.